

Deepfake and Inpainting Detection- A State-of-the-Art Literature Review

The "Industrial Revolution" of Synthetic Media and the Crisis of Trust

- **The Shift to Hyper-Realism** We have witnessed a tectonic shift from the "glitchy" GAN Era to the **Diffusion Era**.
- With the democratization of high-fidelity models like **Stable Diffusion 3**, we are no longer dealing with obvious artifacts but with synthetic media that is "statistically perfect".
- **The "Generalization Gap"** This transition has rendered traditional "Artifact-Centric" detection obsolete.
- A critical gap has emerged: academic models that excel in the lab now fail when tested against "in-the-wild" social media content.
- **Redefining Forensics** To restore reliability, we must pivot away from simple pixel-level classification. The new forensic frontier requires us to analyze **Semantic & Physical Consistency**-looking for logical flaws in the image rather than just noise.

The "Generalization Gap" (Data Analysis)

Quantifying the Failure of Legacy Forensics

The "AUC Collapse":

Analysis of the **DeepFake-Eval-2024** benchmark (published March 2025) reveals a catastrophic performance drop.

State-of-the-art open-source models suffered an average **45–50% decrease in AUC** when tested on social media content (TikTok, X) compared to academic datasets (FaceForensics++).

Specific Degradation Metrics:

- **Video Detectors:** ~50% performance drop.
- **Audio Detectors:** ~48% performance drop.
- **Image Detectors:** ~45% performance drop.

Why?

The "Local Realism" of diffusion models effectively scrubs the high-frequency spectral traces that previous detectors relied upon, while social media compression further masks these signals.

Taxonomy of Manipulation (2024–2026)

Categorizing the Modern Threat Landscape

1

Generative Synthesis (Full Synthesis):

- *Tech:* Shift from U-Net LDMs to **Diffusion Transformers (DiTs)** like FLUX.
- *Trace:* "Spectral roll-off" in high frequencies due to VAE upsampling.

2

Deepfakes (Identity Manipulation):

- *Face Swapping:* One-shot capability (e.g., **Ghost 2.0**, **LivePortrait**) requiring only a single source image.
- *Reenactment:* "Cloned" video avatars driving expression via source video.

3

Text-Guided Inpainting (The New Frontier):

- *Definition:* Masking and regenerating specific regions to alter context (e.g., changing clothes, removing objects).
- *Challenge:* Unlike "splicing," inpainting **harmonizes** lighting and texture, removing edge artifacts.

4

Outpainting:

Extending image boundaries to hallucinate new environments/contexts.

The Inpainting Paradox

Local Realism vs. Global Inconsistency

The Core Challenge:

- **Local Realism:** Diffusion models generate textures (skin pores, fabric) that are statistically indistinguishable from real photos in small patches (64x64 pixels).
- **Global Failure:** The models lack a true understanding of 3D physics and scene logic.

Detection Vectors:

- **Physical Inconsistency:** Mismatched lighting directions, incorrect shadow fall-off, or impossible reflections.
- **VAE Compression Artifacts:** Subtle hue shifts and "smearing" in high-frequency areas (hair, grass) caused by the Latent Diffusion Model's compression step.
- **New Datasets:** Introduction of **SAGI-D** (95k images) to train models specifically on these "hard negatives."

Detection Family 1: Reconstruction & Frequency

The "White-Box" and Signal Processing Approaches

Generative Reconstruction (e.g., End4, DIRE):

- *Hypothesis:* A diffusion model can easily reconstruct images that lie within its own "latent manifold" (Fake = Low Error), but struggles with the entropy of real, complex images (Real = High Error).
- *2025 Advance:* **End4** updates model parameters *during* reconstruction to better align with unseen inpainting models.

Frequency Analysis (e.g., FreqBlender):

- *Mechanism:* Detecting "spectral peaks" or specific noise patterns left by the iterative denoising process.
- *Limitation:* Highly sensitive to **JPEG compression** and resizing (common on social media), which acts as a low-pass filter.

Advanced Cognitive Methodologies (SOTA)

Semantic Reasoning (MLLMs):

- *Models:* **ForgeryGPT, FakeShield.**
- *Innovation:* Treats detection as a **Visual Question Answering (VQA)** task.
- *Output:* Provides natural language explanations (e.g., "The text on the street sign is gibberish") rather than just a binary score.

Multi-Teacher Knowledge Distillation (Re-MTKD):

- *Mechanism:* Uses **Reinforcement Learning (RL)** to aggregate signals from multiple specialized "teacher" networks (splicing, copy-move, inpainting).
- *Goal:* To learn a generalized representation of forgery traces without overfitting to the artifacts of a single generator (like Stable Diffusion).

Comparative Analysis of Leading Models

Benchmarking 2025 Performance

Re-MTKD	Knowledge Distillation	98.1% Accuracy (General benchmarks). Best for broad detection.
InpDiffusion	Generative Masking	97.3% AUC (AutoSplice). Uses diffusion to <i>generate</i> the mask, avoiding jagged edges.
SDiFL	Latent Feature Injection	+12% IoU vs. CAT-Net. Operates in the Latent Space , making it robust to WhatsApp/WeChat compression.
ForgeryGPT	Multimodal LLM	91.4% AUC (IMD2020). High explainability (ROUGE 0.303).

The New Data Ecosystem

Overcoming the "Data Silo" Problem

SAGI-D (2025):

- *Scale:* **95,839 images.**
- *Focus:* "Deceptiveness." Uses Uncertainty-Guided Assessment to include only high-quality fakes. Retraining on this improved generalization by **26.1%.**

So-Fake (2025):

- *Scale:* **Over 2 million images.**
- *Focus:* Social Media Robustness. Collected from Reddit/X to model real-world degradation.

ForensicHub:

A unified benchmark consolidating **23 datasets and 42 baseline models**, allowing instant comparison of new techniques.

Commercial vs. Open-Source Landscape

The Rise of "Ensemble-as-a-Service"

Commercial Strategy:

Moving beyond single models to complex aggregation pipelines.

- **Reality Defender:** Focus on **Enterprise/Government**. Key features: Explainable AI heatmaps and Audio-Visual synchronization checks.
- **TrueMedia.org:** Non-profit aggregator for political/news verification. Widely cited as a "second opinion" layer.
- **Hive:** High accuracy (~82-90%) but vulnerable to specific adversarial "laundering" attacks.

Recommendation:

- **High-Stakes:** Use commercial APIs (Reality Defender) for scale and robustness.
- **Internal/Sensitive:** Deploy open-source models like SDiFL (for localization) and Re-MTKD on local GPUs.

References & Key Literature

Primary Sources (2024–2026)

- **Benchmarks:**

- *DeepFake-Eval-2024: Evaluating Deepfake Detection in the Wild* (2025). [\[link\]](#)
- *ForensicHub: A Unified Benchmark for Digital Forensics* (NeurIPS 2025). [\[link\]](#)

- **Models:**

- *Re-MTKD: Reinforced Multi-Teacher Knowledge Distillation* (AAAI 2025). [\[link\]](#)
- *InpDiffusion: Generative Localization via Conditional Diffusion* (AAAI 2025). [\[link\]](#)
- *SDiFL: Stable Diffusion-driven Framework for Localization* (arXiv 2025). [\[link\]](#)
- *ForgeryGPT: Multimodal LLM for Explainable Detection* (arXiv 2024). [\[link\]](#)

- **Datasets:**

- *SAGI-D: Semantically Aligned and Uncertainty-Guided Inpainting* (ICCV 2025). [\[link\]](#)
- *So-Fake: Social Media Compression and Robustness Dataset* (2025). [\[link\]](#)